Your grade: 100%

Your latest: **100%** • Your highest: **100%** • To pass you need at least 80%. We keep your highest score.

Next item →

1. Suppose you learn a word embedding for a vocabulary of 10000 words. Then the embedding vectors could be 10000 dimensional, so as to capture the full range of variation and meaning in those words. **1 / 1 point**

- ○ True
- ◉ False

Correct
The dimension of word vectors is usually smaller than the size of the vocabulary. Most common sizes for word vectors range between 50 and 1000.

2. True/False: t-SNE is a non-linear dimensionality reduction technique. **1 / 1 point**

- ○ False
- ◉ True

Correct
t-SNE is a non-linear dimensionality reduction technique.

3. Suppose you download a pre-trained word embedding which has been trained on a huge corpus of text. You then use this word embedding to train an RNN for a language task of recognizing if someone is happy from a short snippet of text, using a small training set. **1 / 1 point**

x (input text)	y (happy?)
I'm feeling wonderful today!	1
I'm bummed my cat is ill.	0
Really enjoying this!	1

Then even if the word "ecstatic" does not appear in your small training set, your RNN might reasonably be expected to recognize "I'm ecstatic" as deserving a label $y = 1$.

- ◉ True
- ○ False

Correct
Yes, word vectors empower your model with an incredible ability to generalize. The vector for "ecstatic" would contain a positive/happy connotation which will probably make your model classify the sentence as a "1".

4. Which of these equations do you think should hold for a good word embedding? (Check all that apply) **1 / 1 point**

- ☐ $e_{man} - e_{aunt} \approx e_{woman} - e_{uncle}$
- ☐ $e_{man} - e_{woman} \approx e_{aunt} - e_{uncle}$
- ☑ $e_{man} - e_{uncle} \approx e_{woman} - e_{aunt}$

Correct
The order of words is correct in this analogy.

- ☑ $e_{man} - e_{woman} \approx e_{uncle} - e_{aunt}$

Correct
The order of words is correct in this analogy.

5. Let A be an embedding matrix, and let o_{4567} be a one-hot vector corresponding to word 4567. Then to get the embedding of word 4567, why don't we call $A * o_{4567}$ in Python? **1 / 1 point**

- ○ The correct formula is $A^T * o_{4567}$.
- ◉ It is computationally wasteful.
- ○ None of the answers are correct: calling the Python snippet as described above is fine.
- ○ This doesn't handle unknown words (<UNK>).

Correct
Yes, the element-wise multiplication will be extremely inefficient.

6. When learning word embeddings, we pick a given word and try to predict its surrounding words or vice versa. **1 / 1 point**

- ◉ True
- ○ False

Correct
Word embeddings are learned by picking a given word and trying to predict its surrounding words or vice versa.

7. True/False: In the word2vec algorithm, you estimate *P(t / c)*, where *t* is the target word and *c* is a context word. *t* and *c* are chosen from the training set using *c* as the sequence of all the words in the sentence before *t*. **1 / 1 point**

- ○ True
- ◉ False

Correct
t and *c* are chosen from the training set to be nearby words.

8. Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The word2vec model uses the following softmax function: **1 / 1 point**

$$P(t \mid c) = \frac{e^{\theta_t^T e_c}}{\sum_{t'=1}^{10000} e^{\theta_{t'}^T e_c}}$$

Which of these statements are correct? Check all that apply.

- ☑ θ_t and e_c are both 500 dimensional vectors.

Correct

- ☐ θ_t and e_c are both 10000 dimensional vectors.
- ☑ θ_t and e_c are both trained with an optimization algorithm such as Adam or gradient descent.

Correct

- ☐ After training, we should expect θ_t to be very close to e_c when t and c are the same word.

9. Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings.The GloVe model minimizes this objective: **1 / 1 point**

$$\min \sum_{i=1}^{10,000} \sum_{j=1}^{10,000} f(X_{ij})(\theta_i^T e_j + b_i + b_j' - log X_{ij})^2$$

True/False: θ_i and e_j should be initialized to 0 at the beginning of training.

- ○ True
- ◉ False

Correct
θ_i and e_j should be initialized randomly at the beginning of training.

10. You have trained word embeddings using a text dataset of m_1 words. You are considering using these word embeddings for a language task, for which you have a separate labeled dataset of m_2 words. Keeping in mind that using word embeddings is a form of transfer learning, under which of these circumstances would you expect the word embeddings to be helpful? **1 / 1 point**

- ◉ $m_1 >> m_2$
- ○ $m_1 << m_2$

Correct

Correct